Machine Learning

B. N. M. Institute of Technology

Department: Artificial Intelligence and Machine Learning (AI)

Course: Machine Learning (18AI61) effective from the academic year 2018–2019 under VISVESVARAYA TECHNOLOGICAL UNIVERSITY, BELAGAVI

Instructor: Pradip Kumar Das

Machine Learning Landscape

What is Machine Learning?

Few Examples

Spam Email Filter

Recommender

Optical Character Recognition

Chatbots

Voice Recognition

Facial Recognition

Forecasting

Autonomous Vehicle

Medical Diagnostics

Fraud Detection

and many more...

Some Definitions

"Science of programming computer to learn from data

"field of study that gives computers the ability to learn without being explicitly programmed

*"A computer program is said to learn from experience **E** with respect to some task **T** and some performance measure **P**, if its performance on **T**, as measured by **P**, improves with experience **E***

Why is Machine Learning Required?

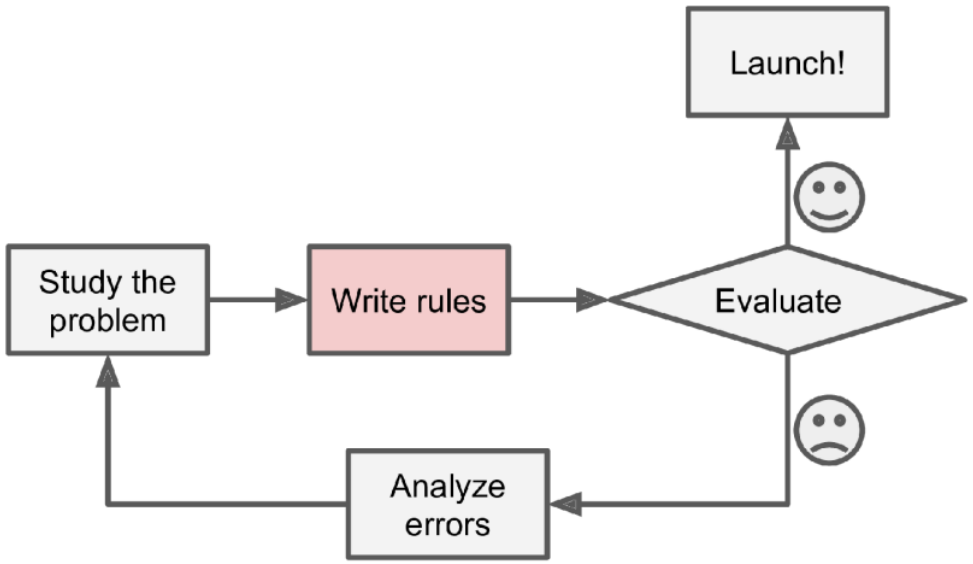

Traditional program writing approach

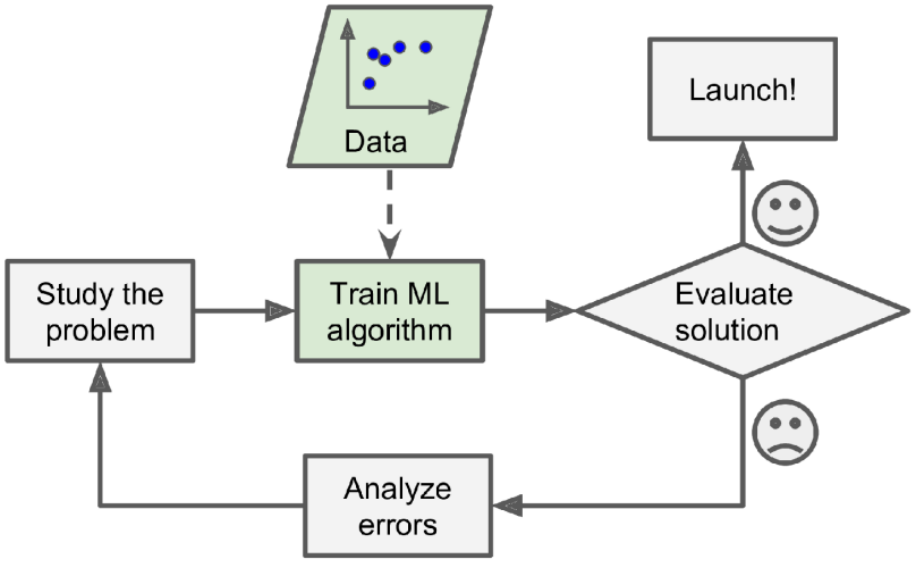

The Machine Learning approach

Why is Machine Learning Required? (Cont.)

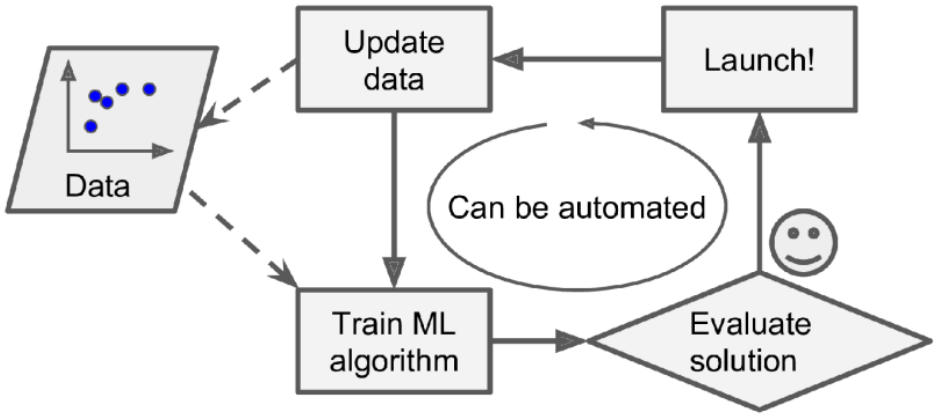

Automatically adapting change

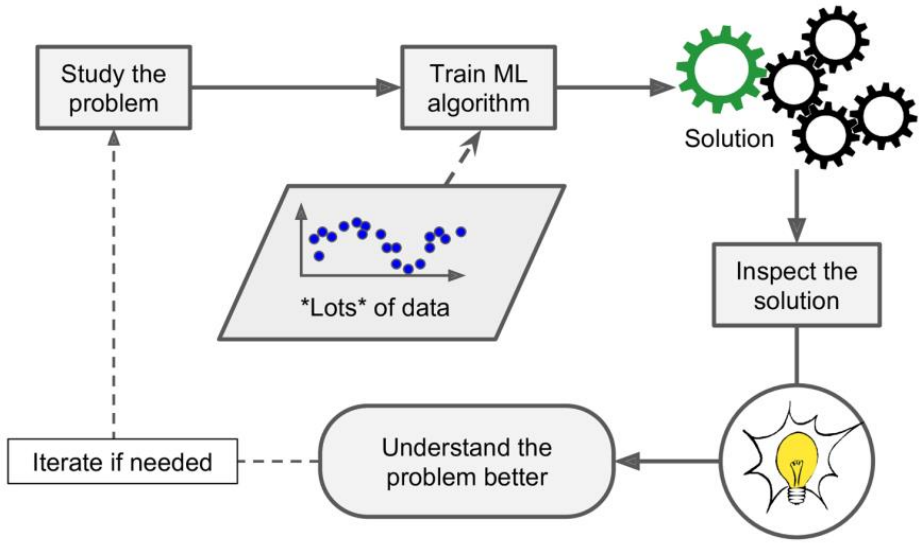

Discovering patterns for human learning

Machine Learning is a Good Option for

- Problems that require a lot of fine-tuning and long list of rules
- Problems that traditional approaches do not yield good results
- Problems where data distribution changes over time
- Getting insights about complex problems from large volume of data

Types of Machine Learning Systems

Whether learning requires human supervision

Types of Machine Learning Systems (Cont.)

Classification & Regression in Supervised Learning

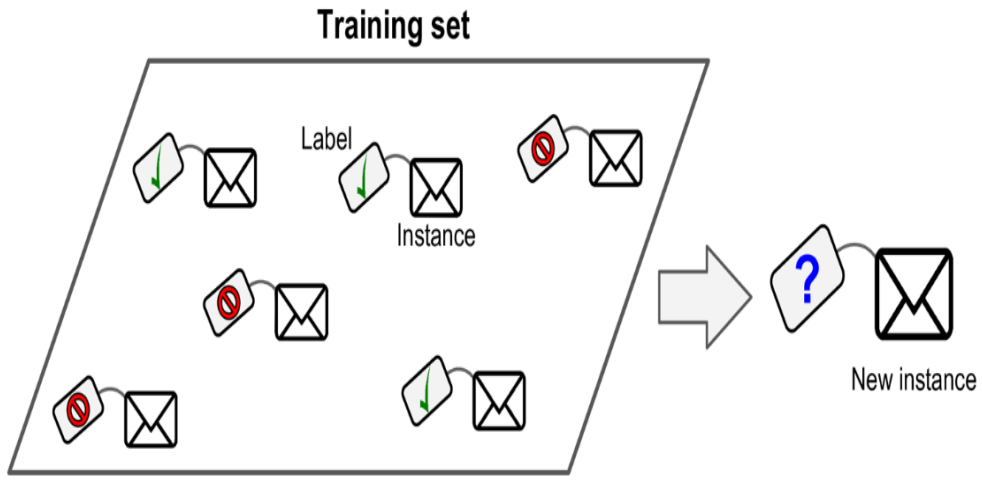

A labeled training set for spam classification

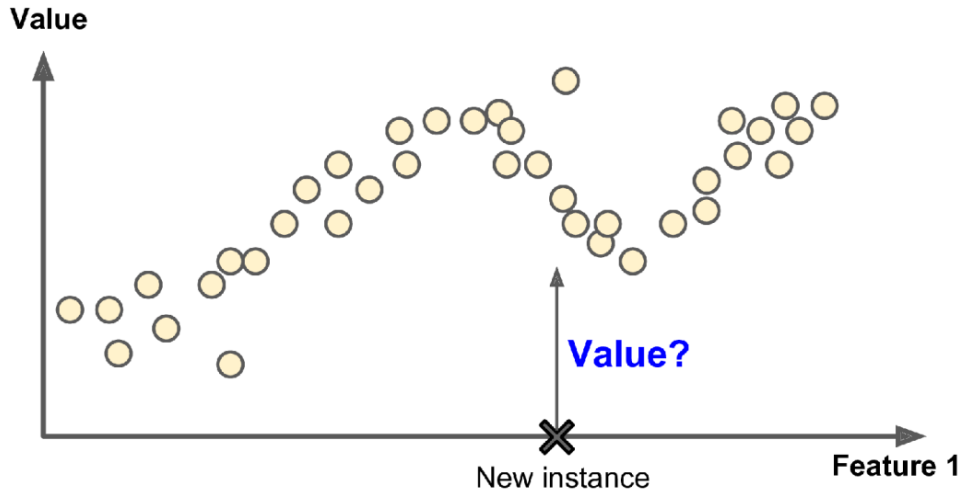

Predicting a continuous value in regression problem

Types of Machine Learning Systems (Cont.)

Clustering in Unsupervised Learning

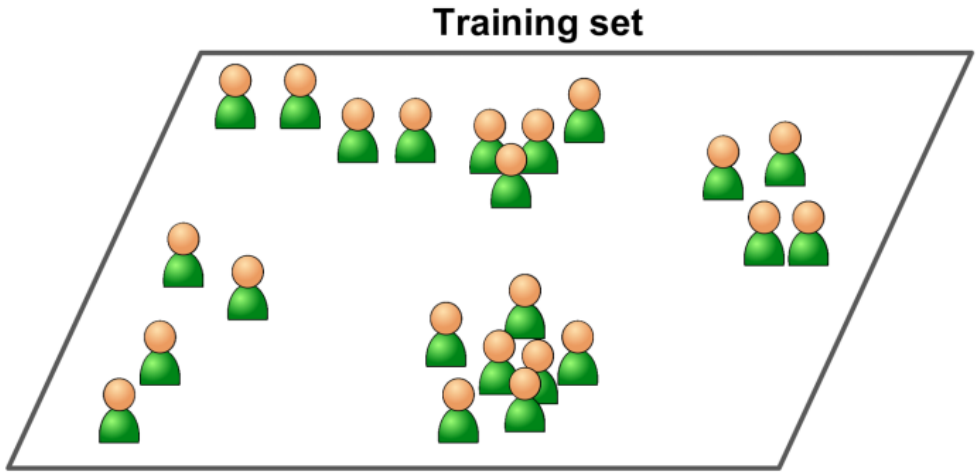

An unlabeled training set for clustering

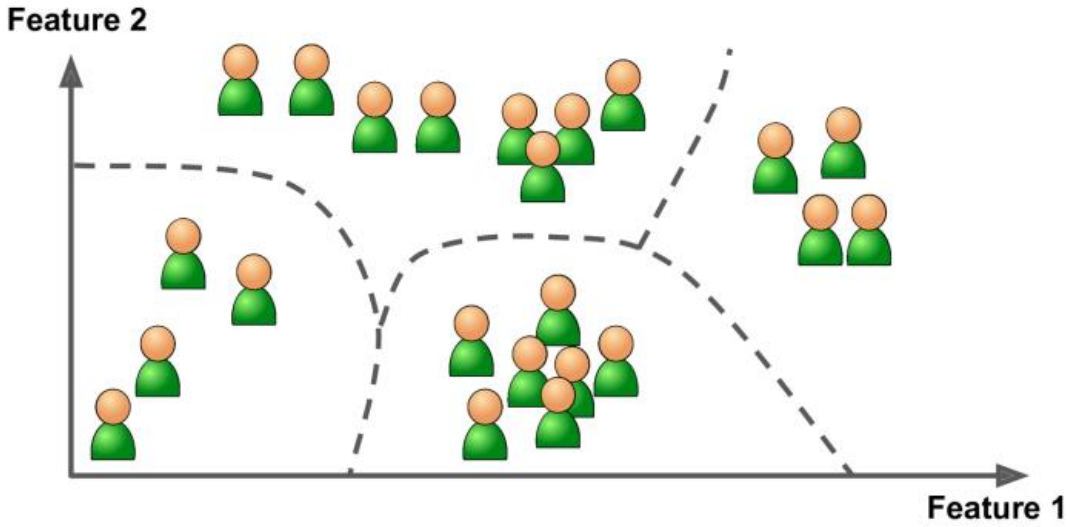

Detecting groups of similar visitor in clustering

Types of Machine Learning Systems (Cont.)

Cluster Visualization in Unsupervised Learning

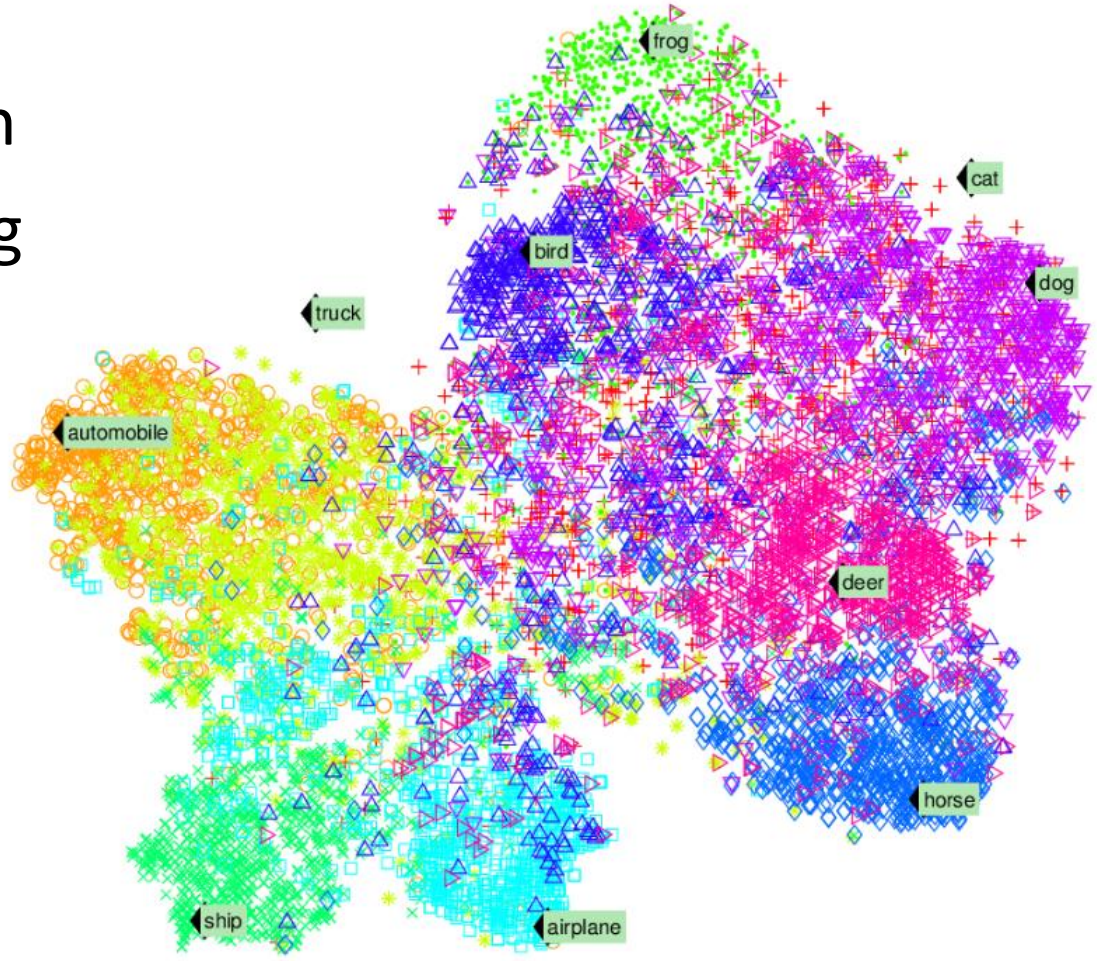

Cluster visualization in unsupervised learning

Types of Machine Learning Systems (Cont.)

Dimensionality Reduction in Unsupervised Learning

- Simplifying data without losing too much information
- Merging correlated features into one
- Feature extraction

Types of Machine Learning Systems (Cont.)

Anomaly and Novelty Detection in Unsupervised Learning

Detection of an anomaly

Types of Machine Learning Systems (Cont.)

Association Rule Learning in Unsupervised Learning

{Potatoes, Onions} => {Burger}

{Barbeque sauce, Potato chips} => {Steak}

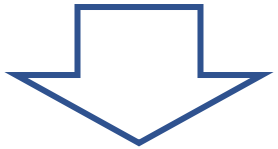

Results in Promotional Pricing and/or Appropriate Product Placement

Sales logs revealing customers' buying pattern

Types of Machine Learning Systems (Cont.)

Classification in Semisupervised Learning

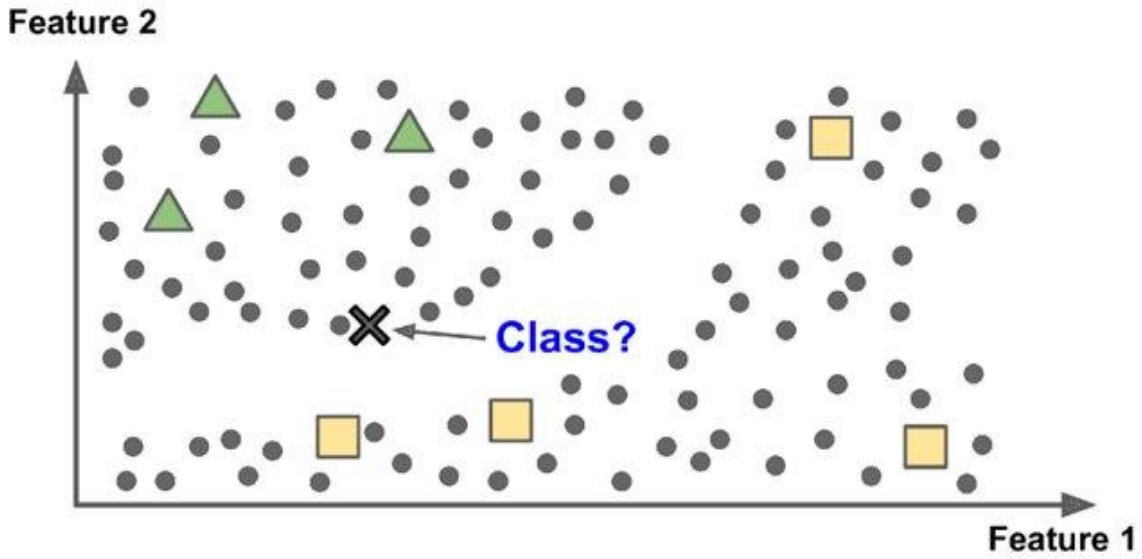

Few labeled examples help classifying new instances

Types of Machine Learning Systems (Cont.)

Building Strategy in Reinforcement Learning

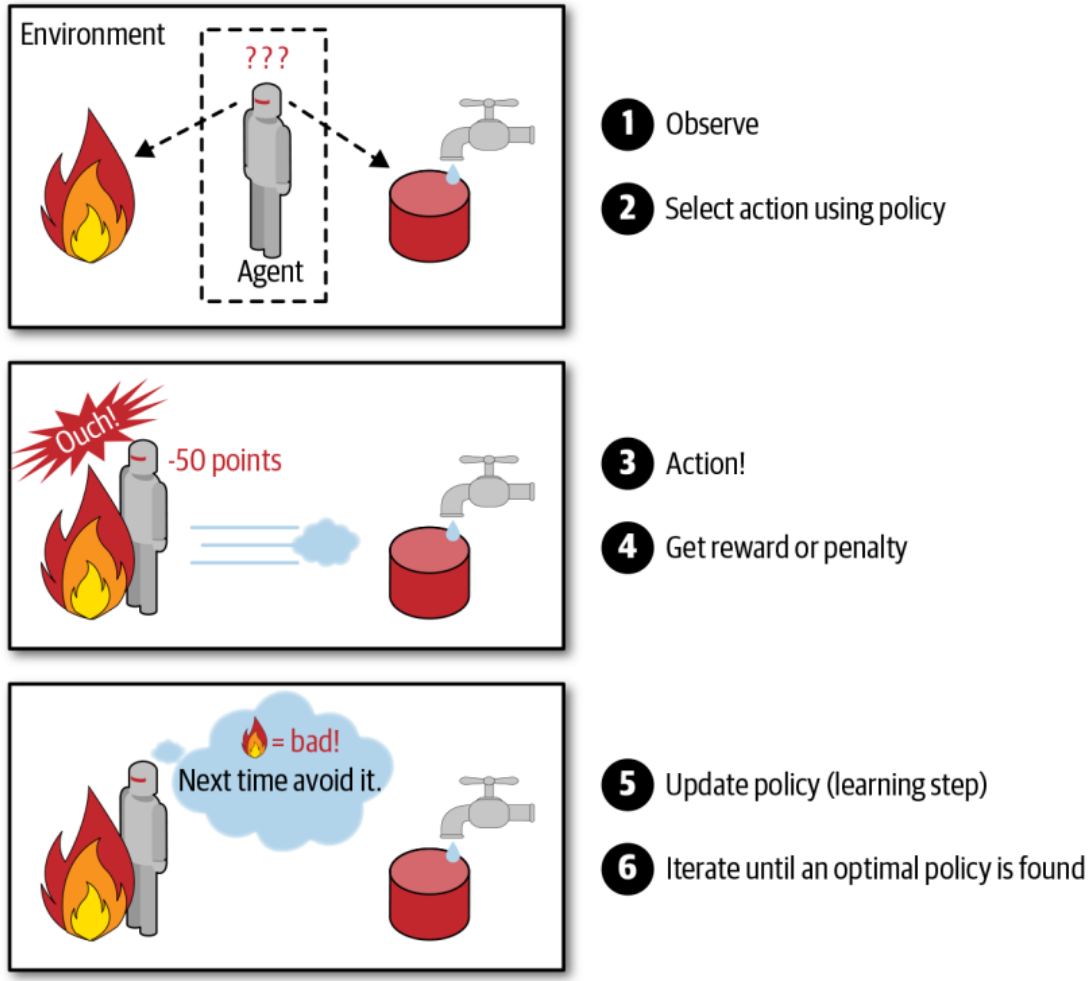

Receiving reward or penalty against its actions and fine-tuning its strategy

Types of Machine Learning Systems (Cont.)

Batch and Online Learning

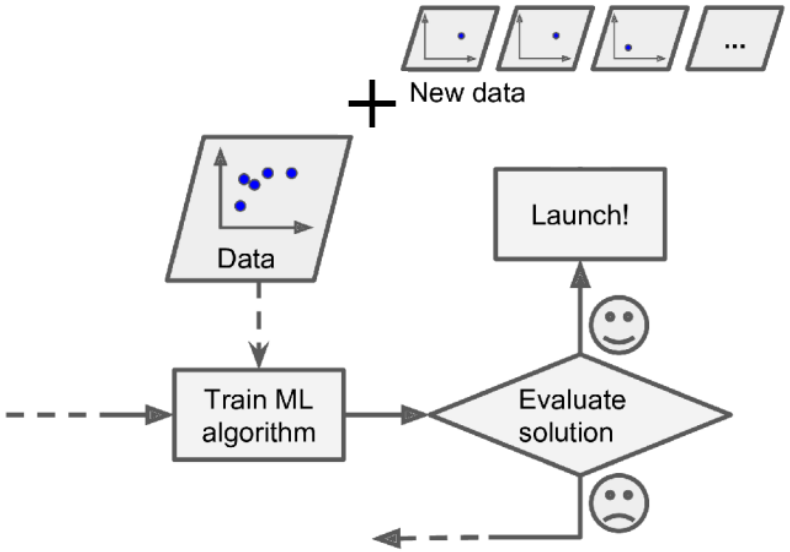

Relearning from scratch on full dataset (old and new data) in Batch Learning

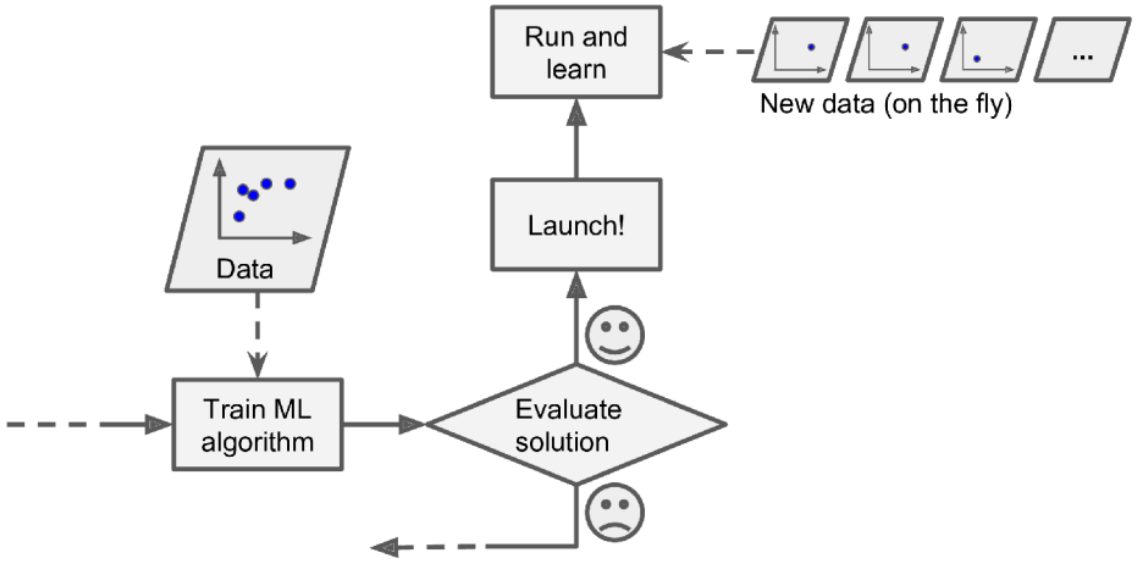

Learning is incremental for new data in Online Learning

Types of Machine Learning Systems (Cont.)

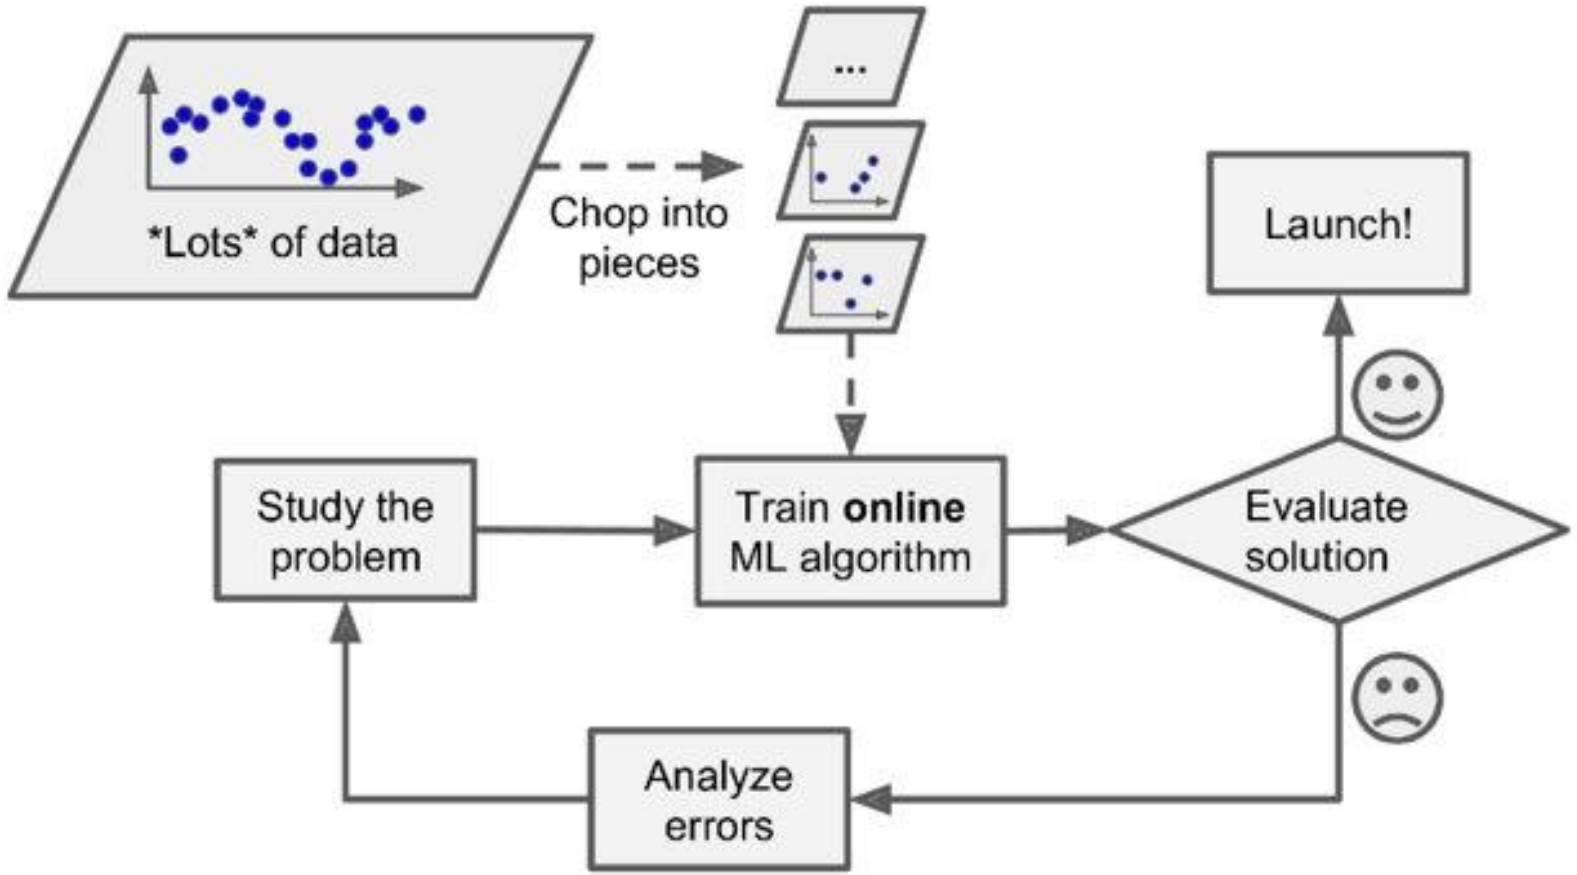

Handling huge dataset in Online Learning

Types of Machine Learning Systems (Cont.)

Instance-based or Model-based Learning systems on how they generalized

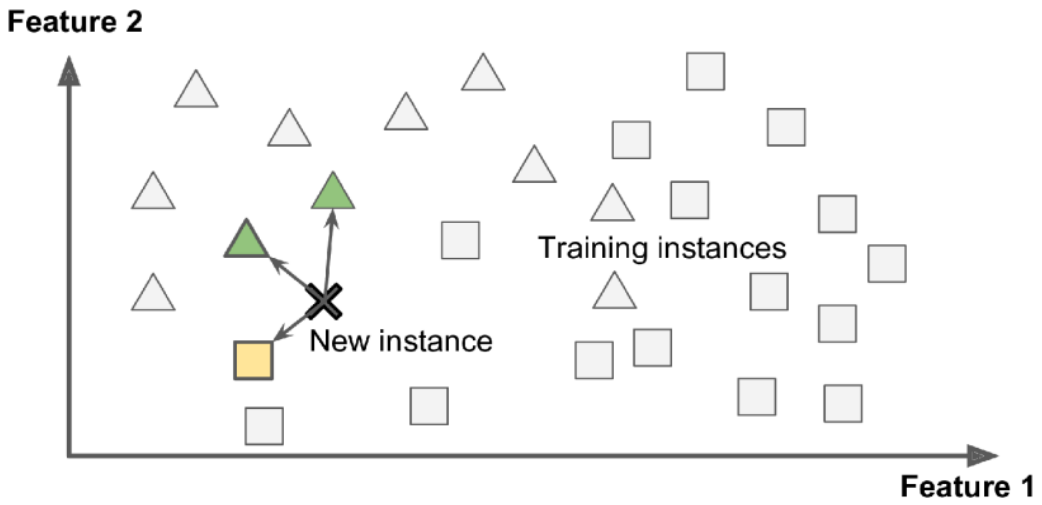

Instance-based Learning is by heart and it then generalizes to new cases by using similarity measures

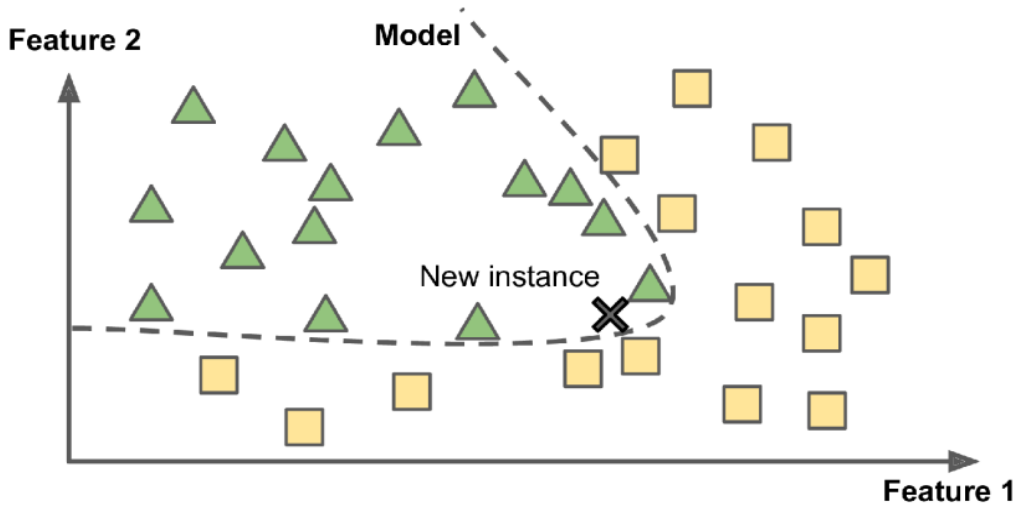

Model is created in Model-based Learning and then it make predictions

Types of Machine Learning Systems (Cont.)

Fitting Model to Data in Model-based Learning

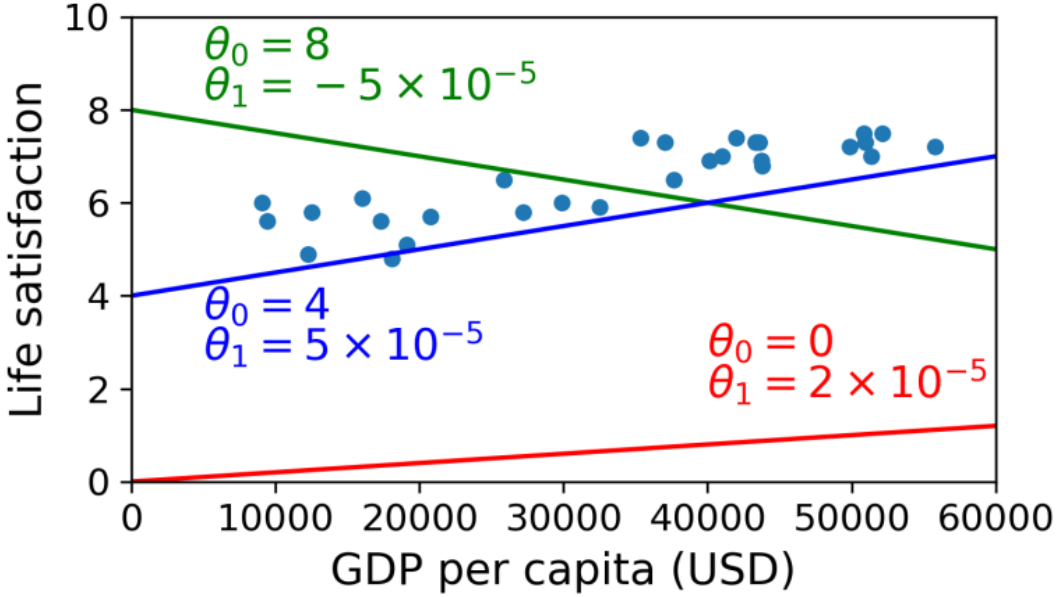

Possible linear models

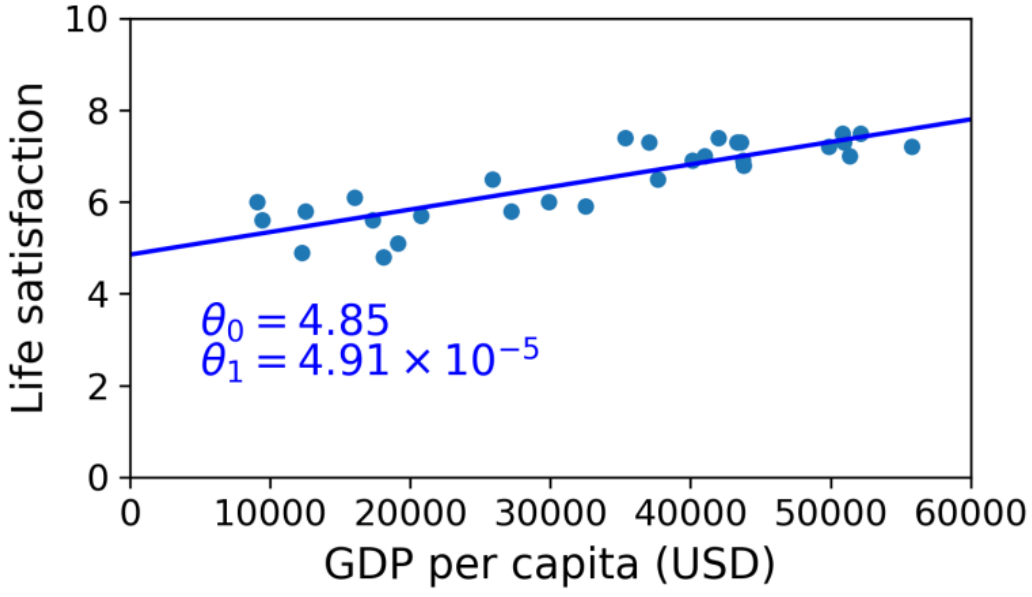

Fitted linear model

Main Challenges in Machine Learning

Insufficient Quantity of Training Data

Trade-off between spending time and money on algorithm development and spending these on corpus development

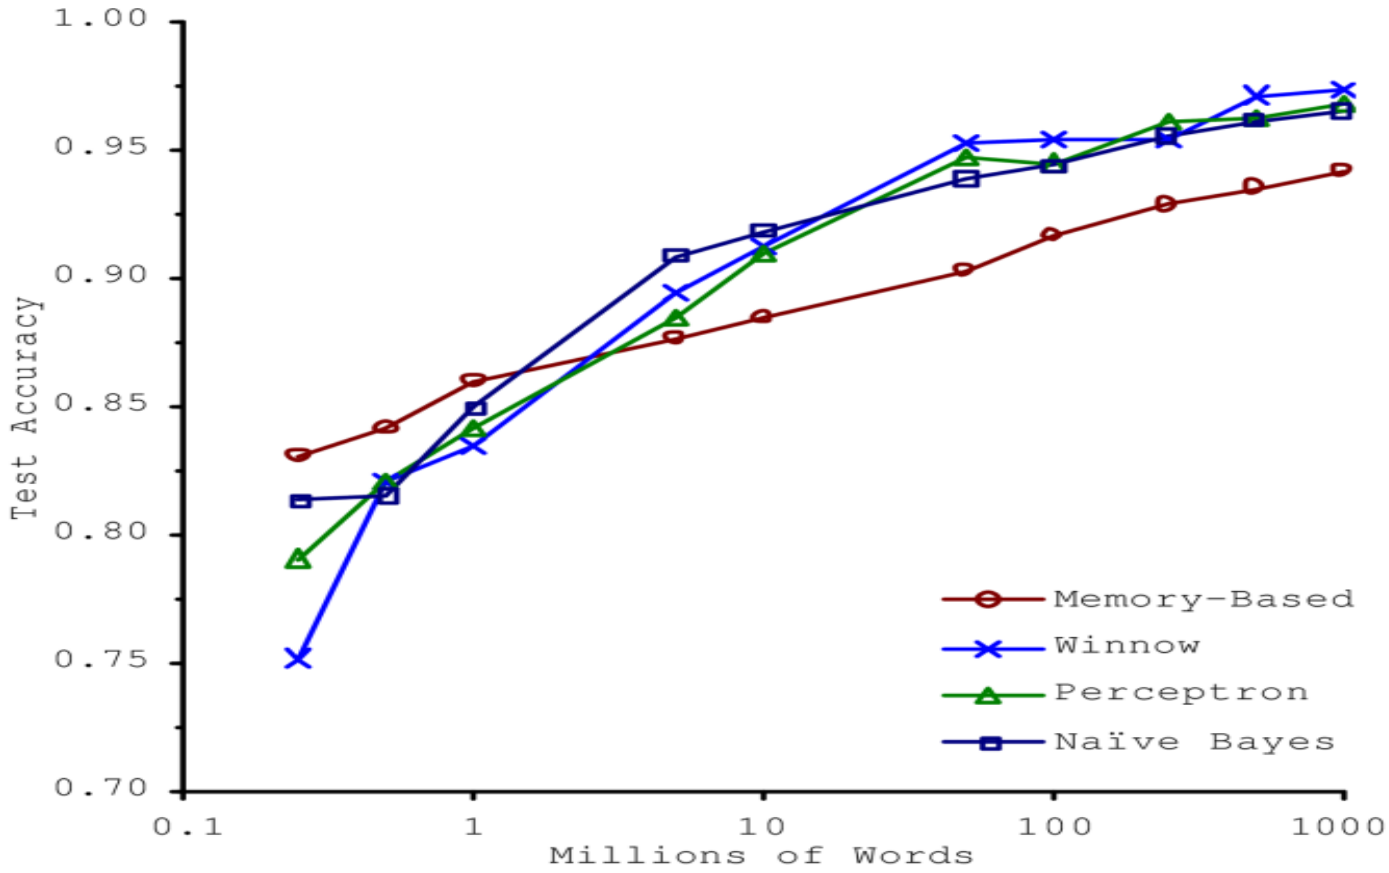

Main Challenges in Machine Learning (Cont.)

Nonrepresentative Training Data

- Sampling Noise
- Sampling Bias

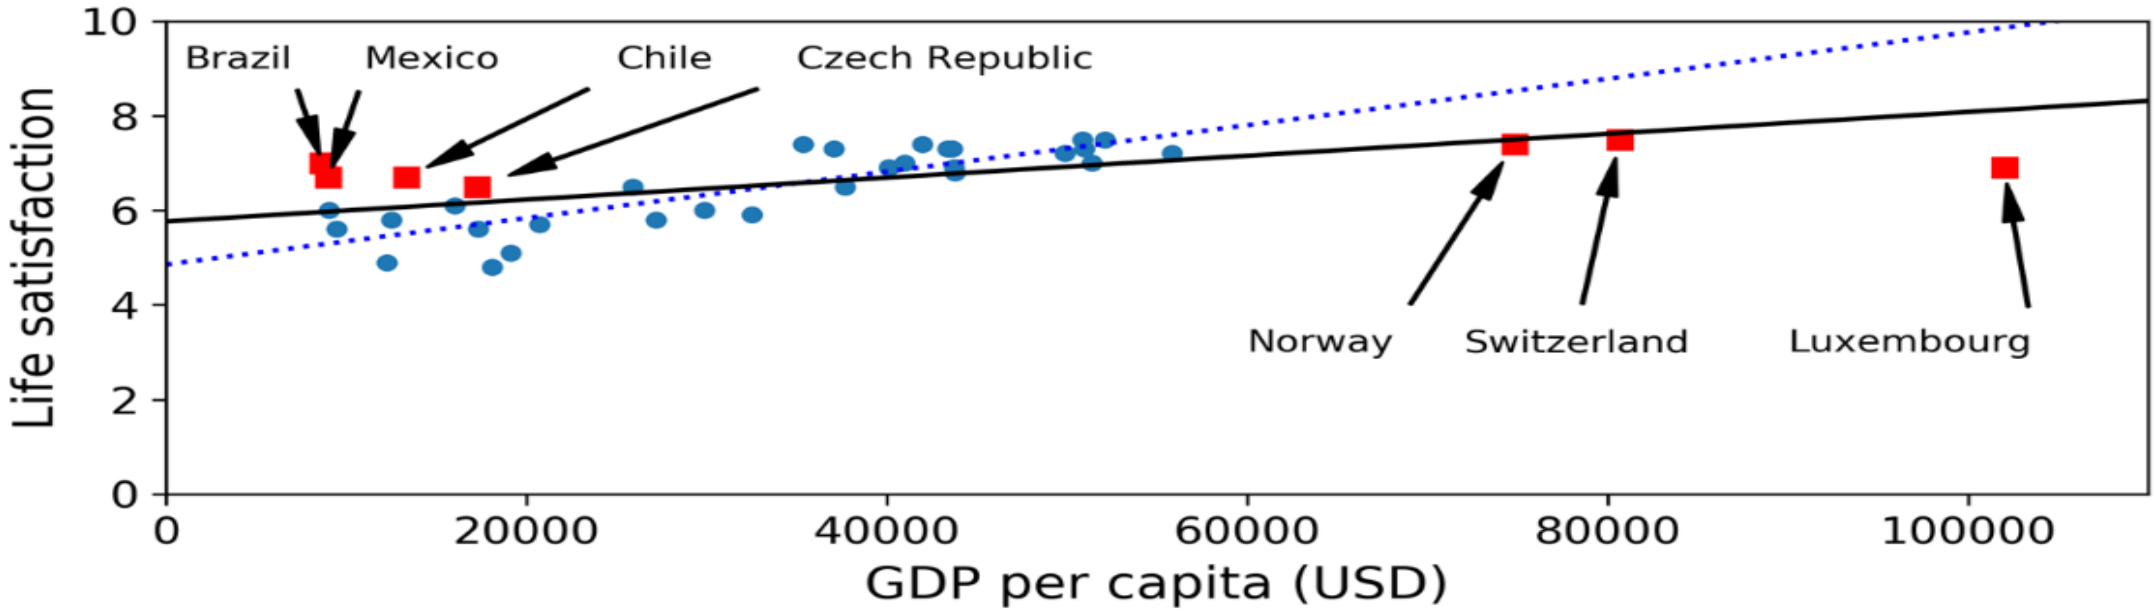

A better model fitted over more representative training data

Main Challenges in Machine Learning (Cont.)

Poor-Quality Data

- Errors
- Missing values
- Outliers

Main Challenges in Machine Learning (Cont.)

Irrelevant Features

Applying Feature Engineering involving the following steps

- Feature selection
- Feature extraction
- Creating new features

Main Challenges in Machine Learning (Cont.)

Overfitting the Training Data

Resolving by

- Simplifying model
- Applying constraints (regularization)
- Gathering more training data
- Fixing data quality issues

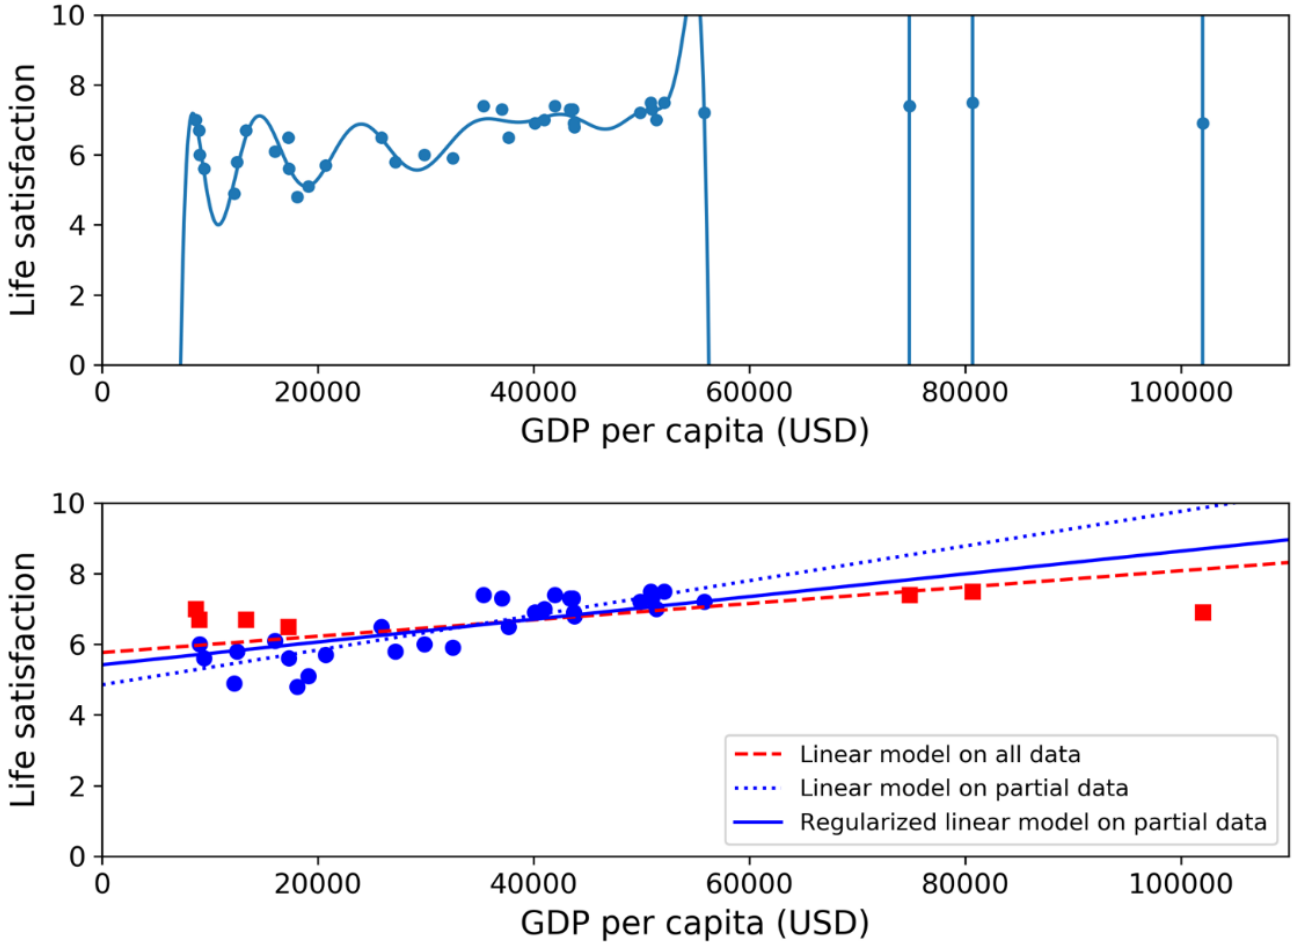

Overfitting and applying regularization to avoid it

Main Challenges in Machine Learning (Cont.)

Underfitting the Training Data

Resolving by

- Selecting more powerful model
- Feeding better features
- Reducing constraints or regularization

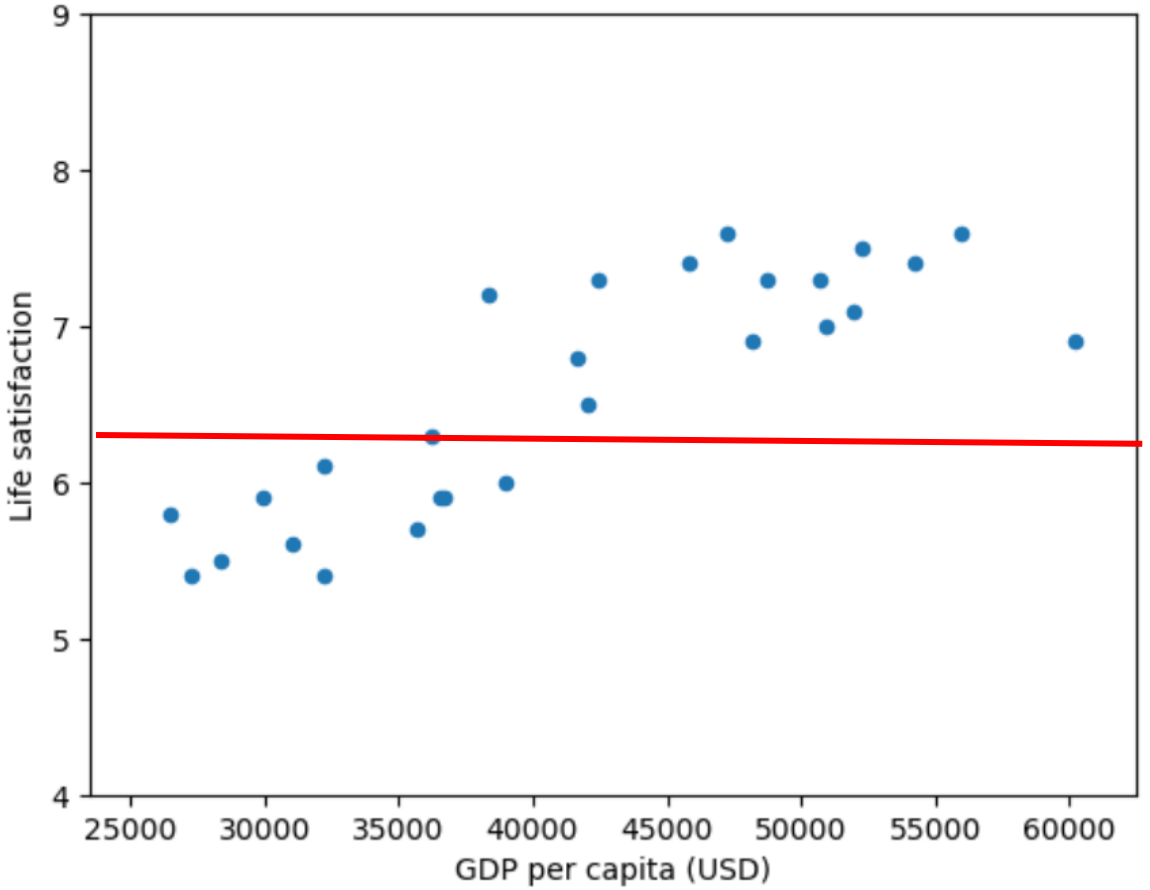

Testing and Validating

Ensuring Model Generalizes Well

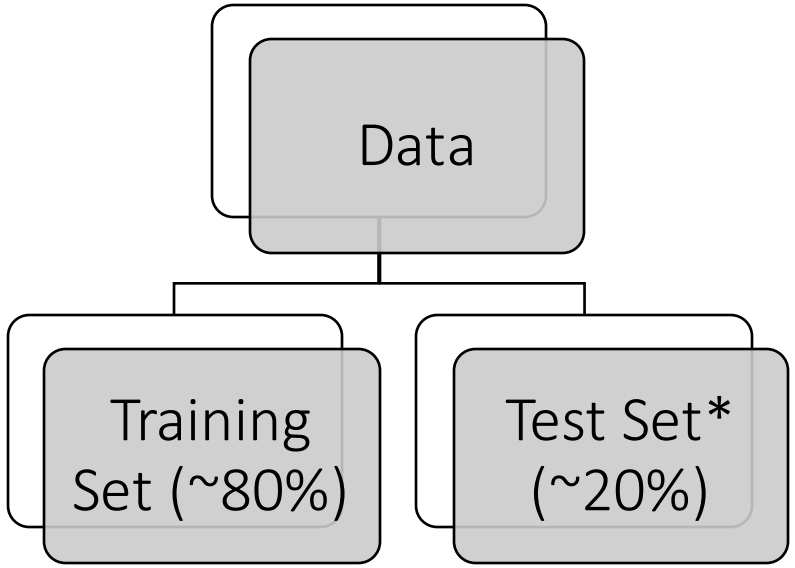

1

Model is built by fitting it on the training data

2

Test data helps estimating how the model generalizes

a.k.a. Hold-out set

A better model fitted over more representative training data

Testing and Validating (Cont.)

- Hyperparameter Tuning
- Model Selection
- Validation
- Cross-validation

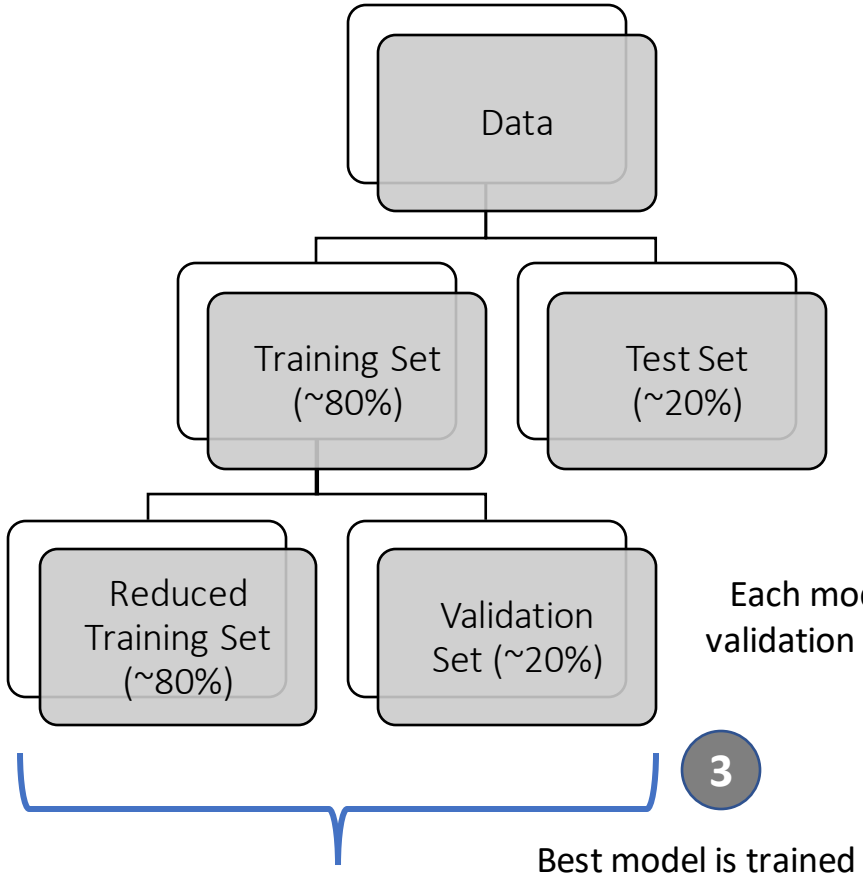

1

Multiple models with different hyperparameters are created on reduced training data set

2

Each models gets tested on validation or hold-out data set

3

Best model is trained on whole training data set

4

Final model gets tested on test set to estimate generalization error

No Free Lunch (NFL) Theorem

- No model that is *a priori* guarantees to be a better one
- Evaluating all models is only way to know which one works best – *Not practical*
- Hence, making *reasonable assumption about data* and *evaluating few reasonable models* is better option

The Pipeline

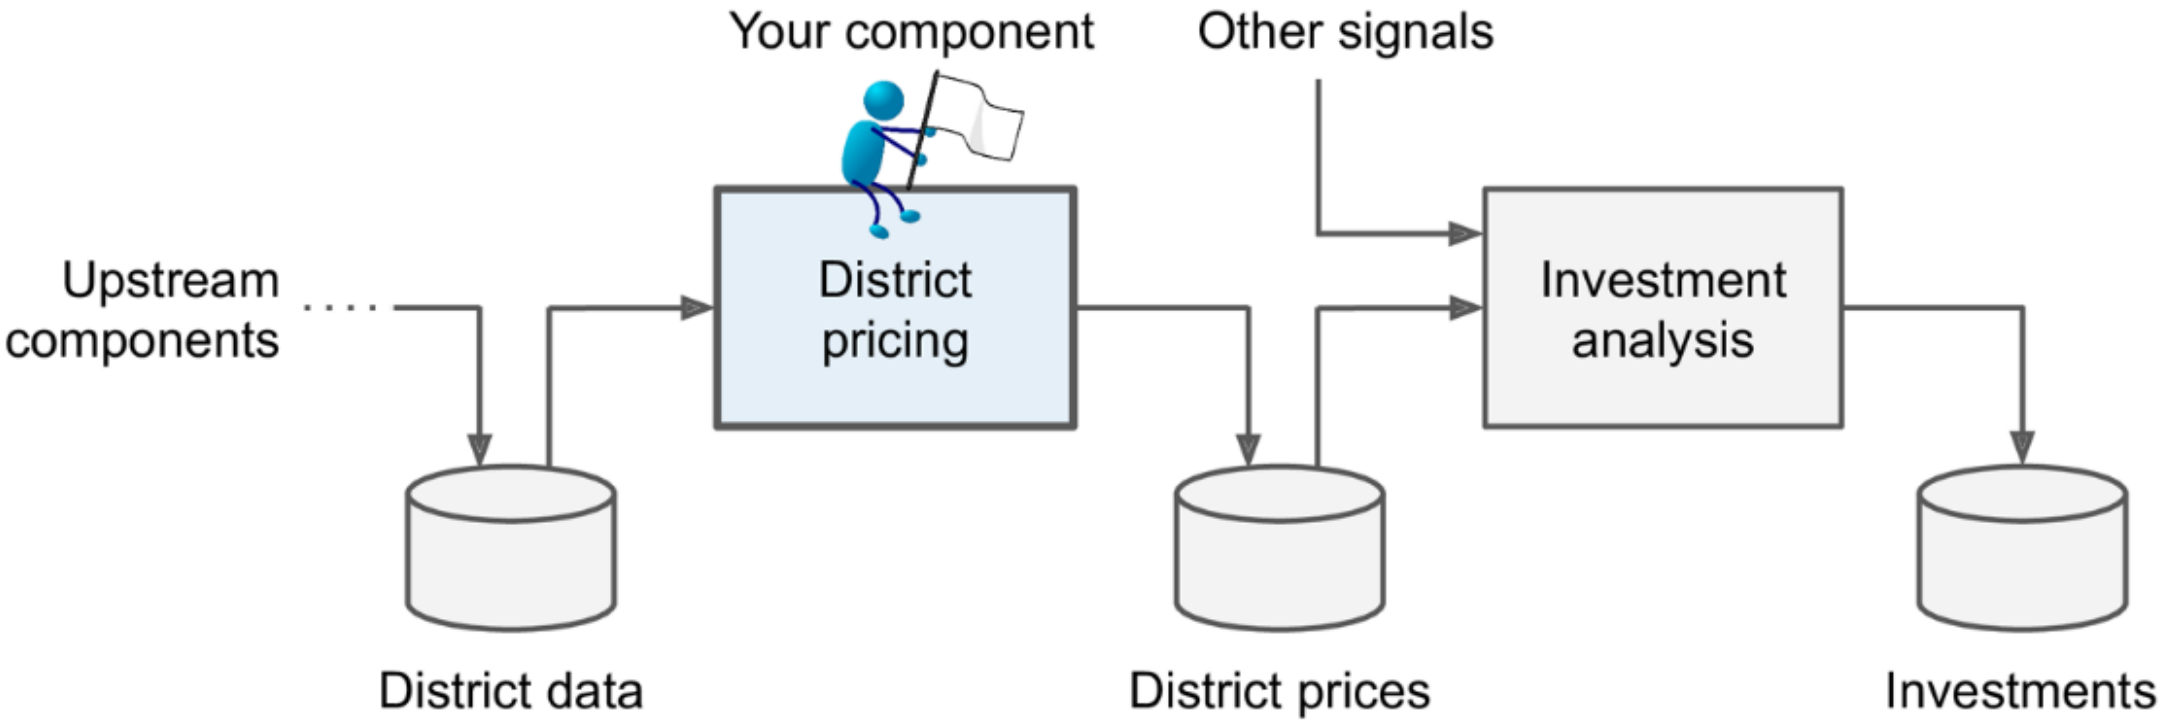

High-level machine learning pipeline for real estate investment

End-to-End Machine Learning Project

Major Steps Involved

1. Looking at the big picture.
2. Getting the data.
3. Discovering and visualizing the data to gain insights.
4. Preparing the data for Machine Learning algorithms.
5. Selecting a model and train it.
6. Fine-tuning your model.
7. Presenting your solution.
8. Launching, monitoring, and maintaining your system.

The Big Picture

- **The Context:** Automatic prediction of district's median house prices using census data
- **Framing the Problem:**
  - Understanding current business problem
    - Time consuming and costly manual operations
    - Performance is not always good
  - Proposed solution
    - A model to predict house price
      - Considering if it should be a supervised, unsupervised or reinforcement learning
      - Considering if it should be classification, (univariate or multivariate) regression or any other type of task
      - Whether it should be batch or online learning
    - Downstream apps to consume the predictions for further business processes and to decide on investments

The Big Picture (Cont.)

- **Deciding on performance metrics:**
  - Means Squared Error (MSE)
  - Root Mean Squared Error (RMSE)
  - Mean Absolute Error (MAE)
- Choosing candidate algorithms
- Realizing development effort

Getting the Data

- Prerequisites (setting up development environment)
- Downloading the relevant data
- Taking a quick look at the data
  - Basic information about dataset
  - Elementary statistics
  - Visualizing
  - Distribution
- Creating test set
  - Randomizing the index of the observations
  - Splitting data into train and test data set (~80:20)
  - Considering stratification, if required

Exploring & Visualizing the Data

Visualizing geographic data

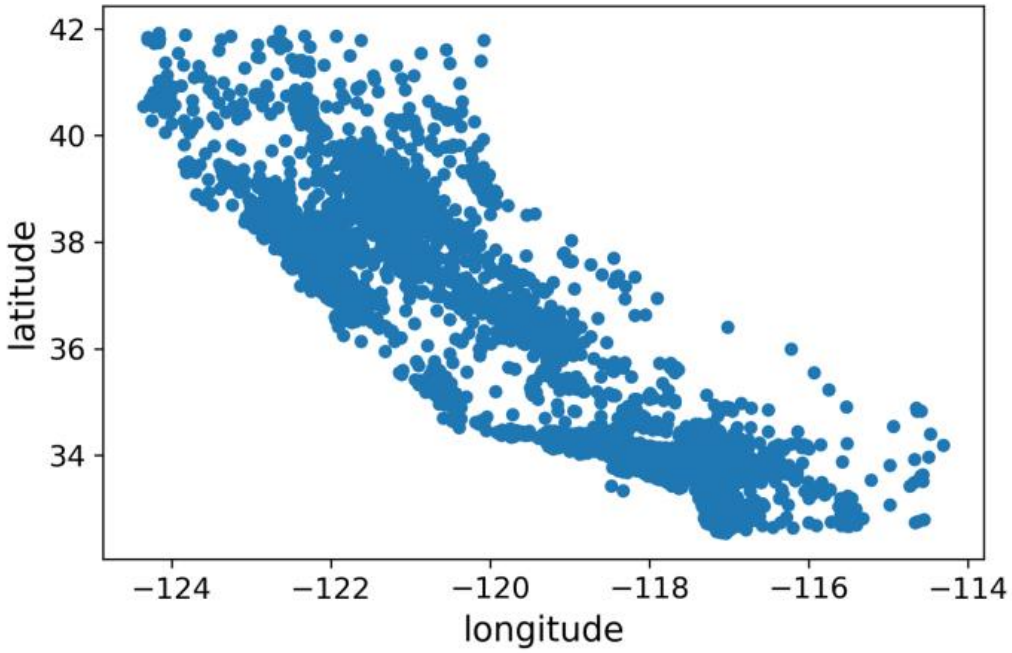

A geographical scatterplot

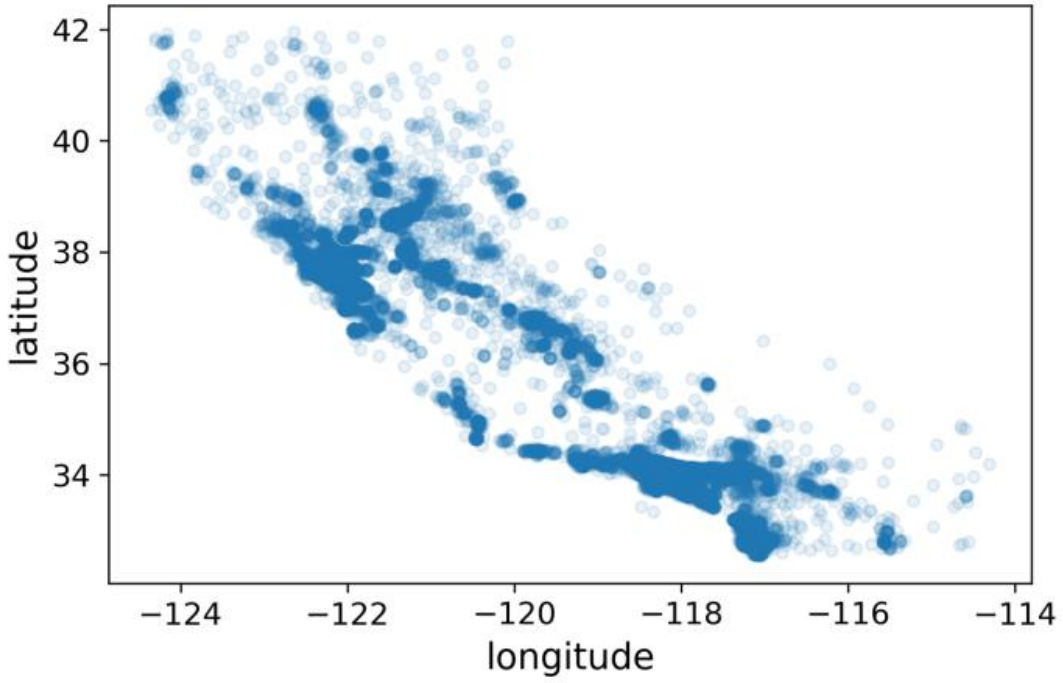

A better visualization of geographical data highlighting high-density areas

Exploring & Visualizing the Data (Cont.)

Better visualization

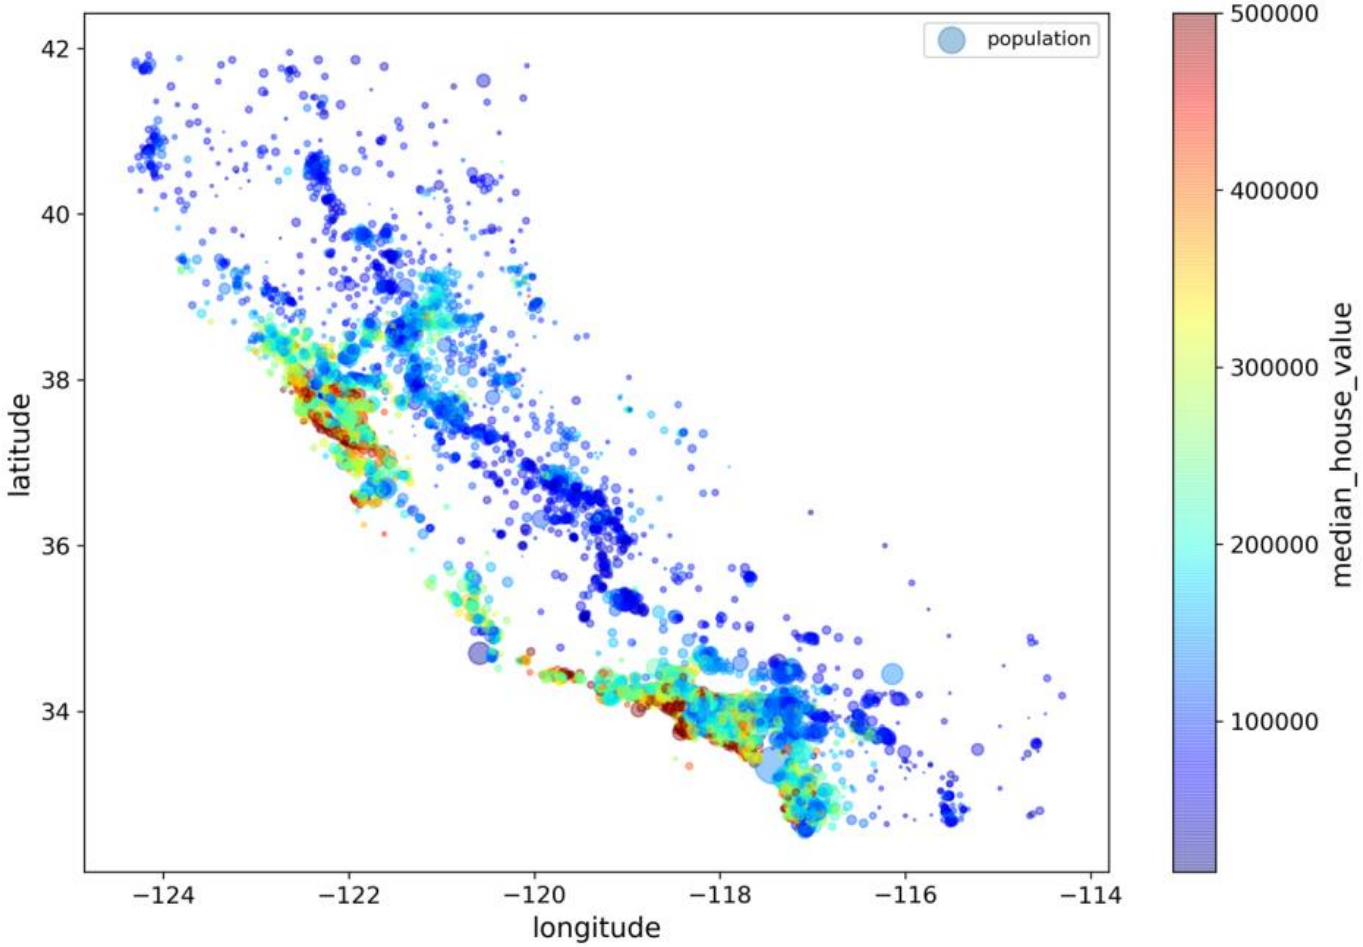

California housing prices: red is expensive, blue is cheap, larger circles indicate areas with a larger population

Exploring & Visualizing the Data (Cont.)

Finding correlations

- Correlation coefficient and its range
- Linear correlation: Positive & Negative correlation
- Nonlinear correlation

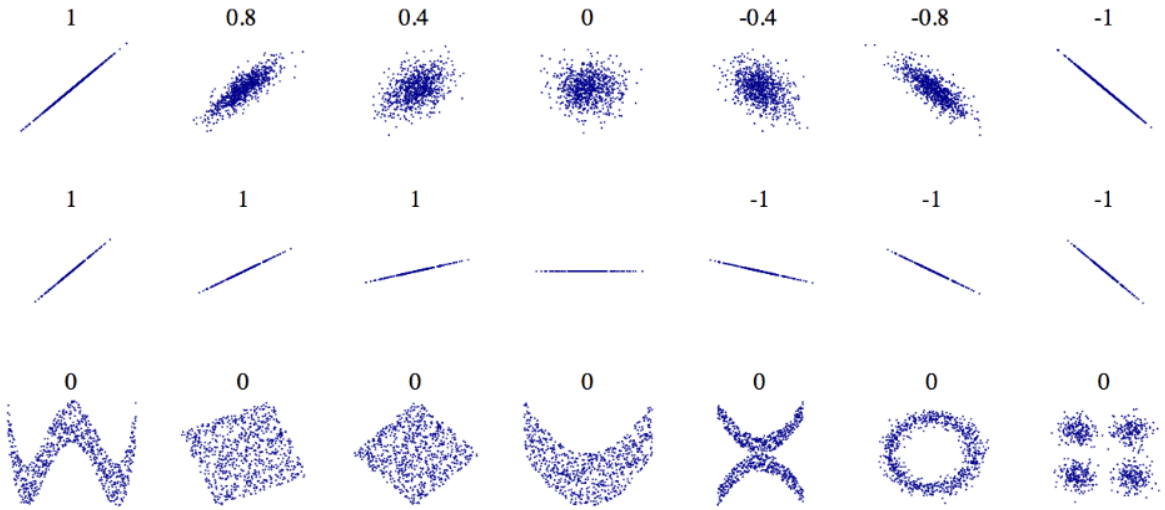

Standard correlation coefficient of various datasets

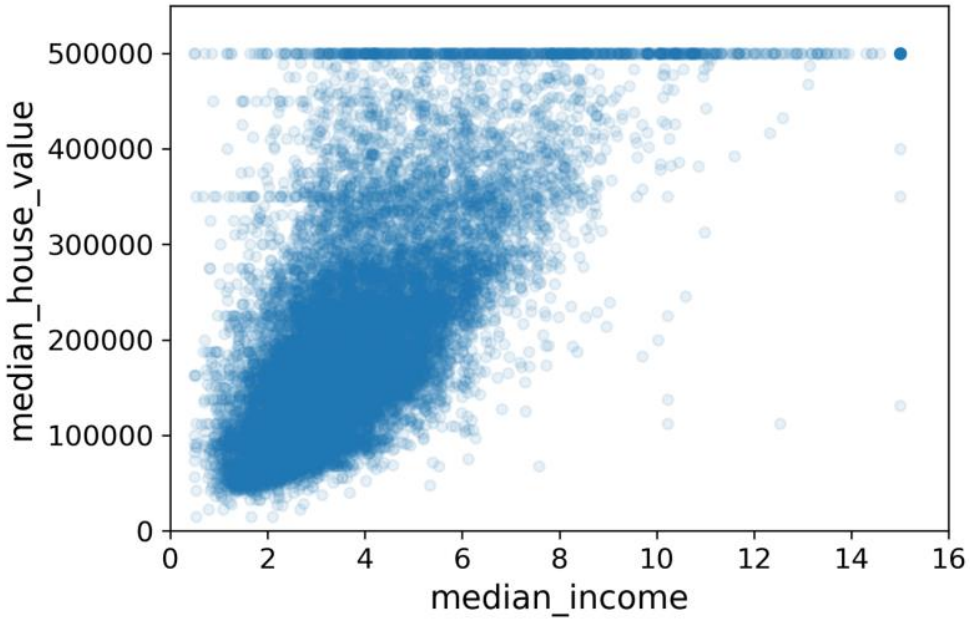

Strong relation between median income and medium house value

Preparing Data

- Data cleaning
  - Handling missing values
    - Removing relevant data, or
    - Removing attribute from entire dataset, or
    - Setting the values with some value (zero, mean, median, etc.)
  - Handling text and categorical attributes
    - Ordinal encoding
    - One-hot encoding
  - Transforming data
    - Scaling features
      - Min-max
      - Standardization
    - Pipelining for transformations

Selecting & Training a Model

- Training model
  - Simple models
  - Complex models
- Evaluating model
  - Evaluation on training set
  - Evaluation over cross-validation
- Considering better model (if applicable)

Fine-tuning Model

- Grid searching
- Randomized searching
- Ensembling
- Analyzing model performances
- Evaluating final model against test set

Launching, Monitoring & Maintaining Models

- Deployment
  - Deploying model as web service
  - Deploying model over Cloud

Model deployed as a web service and being consumed by an application

- Monitoring
  - Monitoring model performance on regular basis
  - Setting up alerts to be triggered when performance falls below threshold
- Maintenance
  - Scripting for automatic model building and hyperparameters tuning
  - Keeping model backup
  - Versioning data sets